

Arquiteturas de Comunicação

Project report

Bruno Gomes, Diogo Silva



Universidade de Aveiro

Arquiteturas de Comunicação

DEPARTAMENTO DE ELETRÓNICA
TELECOMUNICAÇÕES E INFORMÁTICA

Project report

Bruno Gomes, Diogo Silva
(103320) brunofgomes@ua.pt, (104341) diogobranco.as@ua.pt

30/12/2024

Abstract

This project focuses on the technical design, configuration, and testing of a Content Delivery Network (CDN) infrastructure to support a multi-client enterprise environment. The project involves the company CDN4ALL LLC, a provider of Infrastructure-as-a-Service (IaaS) solutions, which operates across geographically distributed Points of Presence (PoPs) in Porto, Lisbon, Barcelona, and Chicago.

Key objectives:

1. **Private Network Design:** Implementing private networks for both large clients and small-to-medium enterprises (SMEs) with guaranteed bandwidth allocation and advanced routing services.
2. **Traffic Differentiation:** Configuring traffic prioritization policies to ensure service quality and efficiency.
3. **Integration of Advanced Networking Technologies:** Utilizing MPLS VPNs, VXLAN, and BGP EVPN mechanisms for robust connectivity and client-specific requirements.

The project leverages VyOS and Cisco C7200 devices for network infrastructure, with specific emphasis on efficient interconnection, advanced service routing, and assured forwarding policies.

Contents

1	Basic assembly and core connectivity	1
1.1	Network diagram	1
1.2	Implementation	2
2	SME1, SME2, and SME3 association private network	6
2.1	Implementation	6
3	SME1, SME2, and SME3 CDN advanced service routing	10
3.1	Implementation	10
4	Client LA Layer 2 private network	16
4.1	Assigned IPs to VPCs	16
4.2	Implementation	16
5	Client LA Layer 2 traffic differentiation	26
6	SME1, SME2, and SME3 private networks traffic reserves	31
7	Conclusion	32

List of Figures

1.1	Planned network topology	1
1.2	Core1 router learned routes	2
1.3	Core1 router MPLS forwarding-table	3
1.4	Lisbon router learned routes	3
1.5	Lisbon router MPLS forwarding-table	3
1.6	Chicago router learned routes	4
1.7	Chicago router MPLS forwarding-table	4
1.8	Porto router learned routes	4
1.9	Porto router MPLS forwarding-table	5
1.10	Barcelona router learned routes	5
1.11	Barcelona router MPLS forwarding-table	5
2.1	SME VRF instance from RL2	6
2.2	RL2 configuration	7
2.3	Attributed label to SME clients	7
2.4	Packet capture that shows ping performed from SME2 to SME1	8
2.5	SME2 to SME1 connectivity	8
2.6	RL2 MPLS forwarding-table	8
2.7	RL2 VRF table	9
2.8	RP2 VRF table	9
2.9	RB1 VRF table	9
3.1	Created views for each SME client	11
3.2	SME1 client database	11
3.3	SME2 client database	12
3.4	SME3 client database	12
3.5	SME1 client query	13
3.6	Master Server response	14
3.7	Successful ping from SME1 client to "cdn4all.com" domain	14
3.8	Successful ping from SME2 client to "cdn4all.com" domain	15
3.9	Successful ping from SME3 client to "cdn4all.com" domain	15
4.1	Network diagram with VPCs	16
4.2	Switch2 configuration	17
4.3	RC1 Bridge 1010 and Bridge 1020	18

4.4	RC1 Bridge 1030 and VIFs	19
4.5	RC1 VXLANs configurations	20
4.6	RC1 BGP neighbors	20
4.7	RC1 L2VPN/EVPN	21
4.8	RC1 OSPF Process	21
4.9	PC4 to PC8 ping	22
4.10	Wireshark capture from VLAN 10 ping	22
4.11	PC1 to PC9 ping	23
4.12	Wireshark capture from VLAN 20 ping	23
4.13	PC6 to PC2 ping	23
4.14	Wireshark capture from VLAN 30 ping	24
4.15	RC1 EVPN prefixes part 1	24
4.16	RC1 EVPN prefixes part 2	25
5.1	Lisbon router class and policy maps	26
5.2	Access-list 101 from Lisbon router	27
5.3	Lisbon router interfaces	27
5.4	Core1 router interfaces	28
5.5	Core1 router class and policy maps	29
5.6	Performed ping from PC1 to PC9	29

Chapter 1

Basic assembly and core connectivity

1.1 Network diagram

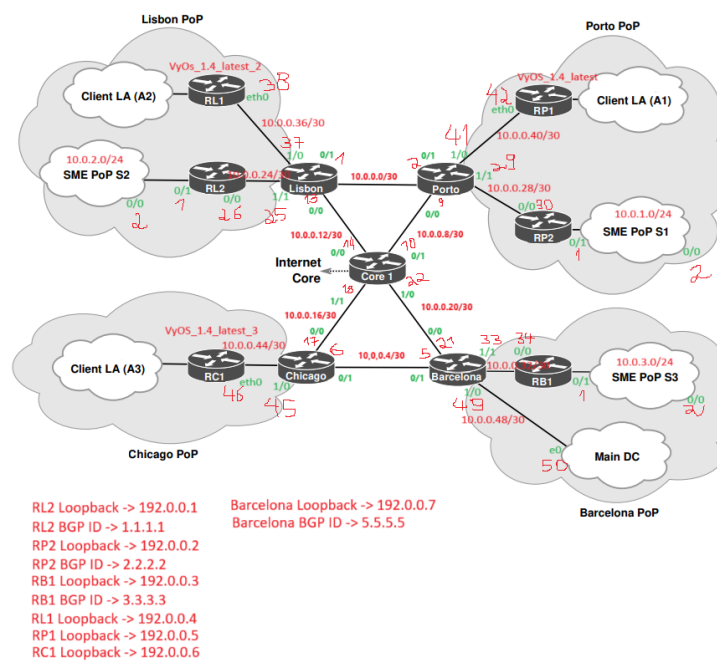


Figure 1.1: Planned network topology

1.2 Implementation

As seen on the network diagram above, we used several IPv4 networks with a mask of /30, we started in the address 10.0.0.0 and then used as many networks necessary to obtain full connectivity, we decided to use a mask of /30 because in this way we do not waste as many addresses if we used a /24 mask for instance.

For the core implementation of this network, we decided to run OSPF and MPLS protocols on routers, Lisbon, Porto, Core1, Chicago and Barcelona. RL2, RP2 and RB1 are also running these protocols but only on the MPLS core side, on the client side there is no OSPF nor MPLS. OSPF was used in order to provide the best path for routing data to its intended destination, and MPLS LDP was used to build label-switch path (LSP) databases that are used to forward traffic through MPLS networks.

MPLS is important for our project, because it allow us to optimize traffic forwarding based on the different labels attributed to the routing paths, this aspect works well with traffic engineering methods which should have been explored in this project.

Instead of performing pings from every single point in the network to verify full connectivity, we prefer to demonstrate the routes learned by OSPF on each main router, along with the MPLS forwarding table for each router.

```
10.0.0.0/8 is variably subnetted, 17 subnets, 2 masks
O    10.0.0.0/30 [110/2] via 10.0.0.13, 00:15:19, FastEthernet0/0
    [110/2] via 10.0.0.9, 00:15:19, FastEthernet0/1
O    10.0.0.4/30 [110/2] via 10.0.0.21, 00:15:19, FastEthernet1/0
    [110/2] via 10.0.0.17, 00:15:19, FastEthernet1/1
C    10.0.0.8/30 is directly connected, FastEthernet0/1
L    10.0.0.10/32 is directly connected, FastEthernet0/1
C    10.0.0.12/30 is directly connected, FastEthernet0/0
L    10.0.0.14/32 is directly connected, FastEthernet0/0
C    10.0.0.16/30 is directly connected, FastEthernet1/1
L    10.0.0.18/32 is directly connected, FastEthernet1/1
C    10.0.0.20/30 is directly connected, FastEthernet1/0
L    10.0.0.22/32 is directly connected, FastEthernet1/0
O    10.0.0.24/30 [110/2] via 10.0.0.13, 00:15:30, FastEthernet0/0
O    10.0.0.28/30 [110/2] via 10.0.0.9, 00:15:20, FastEthernet0/1
O    10.0.0.32/30 [110/2] via 10.0.0.21, 00:15:20, FastEthernet1/0
O    10.0.0.36/30 [110/2] via 10.0.0.13, 00:15:30, FastEthernet0/0
O    10.0.0.40/30 [110/2] via 10.0.0.9, 00:15:20, FastEthernet0/1
O    10.0.0.44/30 [110/2] via 10.0.0.17, 00:15:20, FastEthernet1/1
O    10.0.0.48/30 [110/2] via 10.0.0.21, 00:15:20, FastEthernet1/0
Core1#
```

Figure 1.2: Core1 router learned routes


```
Core1#show mpls forwarding-table
Local   Outgoing Prefix      Bytes Label   Outgoing  Next Hop
Label   Label    or Tunnel Id  Switched      interface
16      Pop Label 10.0.0.36/30  0             Fa0/0      10.0.0.13
17      Pop Label 10.0.0.24/30  0             Fa0/0      10.0.0.13
18      Pop Label 10.0.0.0/30   0             Fa0/1      10.0.0.9
        Pop Label 10.0.0.0/30   0             Fa0/0      10.0.0.13
19      Pop Label 10.0.0.48/30  0             Fa1/0      10.0.0.21
20      Pop Label 10.0.0.44/30  0             Fa1/1      10.0.0.17
21      Pop Label 10.0.0.40/30  0             Fa0/1      10.0.0.9
22      Pop Label 10.0.0.32/30  0             Fa1/0      10.0.0.21
23      Pop Label 10.0.0.28/30  0             Fa0/1      10.0.0.9
24      Pop Label 10.0.0.4/30  0             Fa1/1      10.0.0.17
        Pop Label 10.0.0.4/30  0             Fa1/0      10.0.0.21
```

Figure 1.3: Core1 router MPLS forwarding-table

```
10.0.0.0/8 is variably subnetted, 17 subnets, 2 masks
C    10.0.0.0/30 is directly connected, FastEthernet0/1
L    10.0.0.1/32 is directly connected, FastEthernet0/1
O    10.0.0.4/30 [110/3] via 10.0.0.14, 02:03:48, FastEthernet0/0
O    10.0.0.8/30 [110/2] via 10.0.0.14, 02:03:58, FastEthernet0/0
        [110/2] via 10.0.0.2, 02:03:58, FastEthernet0/1
C    10.0.0.12/30 is directly connected, FastEthernet0/0
L    10.0.0.13/32 is directly connected, FastEthernet0/0
O    10.0.0.16/30 [110/2] via 10.0.0.14, 02:03:58, FastEthernet0/0
O    10.0.0.20/30 [110/2] via 10.0.0.14, 02:03:58, FastEthernet0/0
C    10.0.0.24/30 is directly connected, FastEthernet1/1
L    10.0.0.25/32 is directly connected, FastEthernet1/1
O    10.0.0.28/30 [110/2] via 10.0.0.2, 02:03:58, FastEthernet0/1
O    10.0.0.32/30 [110/3] via 10.0.0.14, 02:03:58, FastEthernet0/0
C    10.0.0.36/30 is directly connected, FastEthernet1/0
L    10.0.0.37/32 is directly connected, FastEthernet1/0
O    10.0.0.40/30 [110/2] via 10.0.0.2, 02:03:58, FastEthernet0/1
O    10.0.0.44/30 [110/3] via 10.0.0.14, 02:03:58, FastEthernet0/0
O    10.0.0.48/30 [110/3] via 10.0.0.14, 02:03:58, FastEthernet0/0
Lisbon#
```

Figure 1.4: Lisbon router learned routes

```
Lisbon#show mpls forwarding-table
Local   Outgoing Prefix      Bytes Label   Outgoing  Next Hop
Label   Label    or Tunnel Id  Switched      interface
16      Pop Label 10.0.0.20/30  0             Fa0/0      10.0.0.14
17      Pop Label 10.0.0.16/30  0             Fa0/0      10.0.0.14
18      Pop Label 10.0.0.8/30  0             Fa0/1      10.0.0.2
        Pop Label 10.0.0.8/30  0             Fa0/0      10.0.0.14
19      19        10.0.0.48/30  0             Fa0/0      10.0.0.14
20      20        10.0.0.44/30  0             Fa0/0      10.0.0.14
21      Pop Label 10.0.0.40/30  0             Fa0/1      10.0.0.2
22      22        10.0.0.32/30  0             Fa0/0      10.0.0.14
23      Pop Label 10.0.0.28/30  0             Fa0/1      10.0.0.2
24      24        10.0.0.4/30   0             Fa0/0      10.0.0.14
```

Figure 1.5: Lisbon router MPLS forwarding-table

```

10.0.0.0/8 is variably subnetted, 16 subnets, 2 masks
O    10.0.0.0/30 [110/3] via 10.0.0.18, 02:22:40, FastEthernet0/0
C    10.0.0.4/30 is directly connected, FastEthernet0/1
L    10.0.0.6/32 is directly connected, FastEthernet0/1
O    10.0.0.8/30 [110/2] via 10.0.0.18, 02:22:40, FastEthernet0/0
O    10.0.0.12/30 [110/2] via 10.0.0.18, 02:22:40, FastEthernet0/0
C    10.0.0.16/30 is directly connected, FastEthernet0/0
L    10.0.0.17/32 is directly connected, FastEthernet0/0
O    10.0.0.20/30 [110/2] via 10.0.0.18, 02:22:40, FastEthernet0/0
    [110/2] via 10.0.0.5, 02:22:30, FastEthernet0/1
O    10.0.0.24/30 [110/3] via 10.0.0.18, 02:22:40, FastEthernet0/0
O    10.0.0.28/30 [110/3] via 10.0.0.18, 02:22:40, FastEthernet0/0
O    10.0.0.32/30 [110/2] via 10.0.0.5, 02:22:30, FastEthernet0/1
O    10.0.0.36/30 [110/3] via 10.0.0.18, 02:22:40, FastEthernet0/0
O    10.0.0.40/30 [110/3] via 10.0.0.18, 02:22:40, FastEthernet0/0
C    10.0.0.44/30 is directly connected, FastEthernet1/0
L    10.0.0.45/32 is directly connected, FastEthernet1/0
O    10.0.0.48/30 [110/2] via 10.0.0.5, 02:22:30, FastEthernet0/1
Chicago#

```

Figure 1.6: Chicago router learned routes

```

Chicago#show mpls forwarding-table
Local   Outgoing   Prefix      Bytes Label   Outgoing   Next Hop
Label   Label      or Tunnel Id Switched      interface
16      Pop Label  10.0.0.48/30 0             Fa0/1       10.0.0.5
17      21         10.0.0.40/30 0             Fa0/0       10.0.0.18
18      16         10.0.0.36/30 0             Fa0/0       10.0.0.18
19      Pop Label  10.0.0.32/30 0             Fa0/1       10.0.0.5
20      23         10.0.0.28/30 0             Fa0/0       10.0.0.18
21      17         10.0.0.24/30 0             Fa0/0       10.0.0.18
22      18         10.0.0.0/30  0             Fa0/0       10.0.0.18
23      Pop Label  10.0.0.12/30 0             Fa0/0       10.0.0.18
24      Pop Label  10.0.0.8/30  0             Fa0/0       10.0.0.18
25      Pop Label  10.0.0.20/30 0             Fa0/1       10.0.0.5
        Pop Label  10.0.0.20/30 0             Fa0/0       10.0.0.18

```

Figure 1.7: Chicago router MPLS forwarding-table

```

10.0.0.0/8 is variably subnetted, 17 subnets, 2 masks
C    10.0.0.0/30 is directly connected, FastEthernet0/1
L    10.0.0.2/32 is directly connected, FastEthernet0/1
O    10.0.0.4/30 [110/3] via 10.0.0.10, 02:28:46, FastEthernet0/0
C    10.0.0.8/30 is directly connected, FastEthernet0/0
L    10.0.0.9/32 is directly connected, FastEthernet0/0
O    10.0.0.12/30 [110/2] via 10.0.0.10, 02:28:56, FastEthernet0/0
    [110/2] via 10.0.0.1, 02:28:56, FastEthernet0/1
O    10.0.0.16/30 [110/2] via 10.0.0.10, 02:28:46, FastEthernet0/0
O    10.0.0.20/30 [110/2] via 10.0.0.10, 02:28:46, FastEthernet0/0
O    10.0.0.24/30 [110/2] via 10.0.0.1, 02:28:56, FastEthernet0/1
C    10.0.0.28/30 is directly connected, FastEthernet1/1
L    10.0.0.29/32 is directly connected, FastEthernet1/1
O    10.0.0.32/30 [110/3] via 10.0.0.10, 02:28:46, FastEthernet0/0
O    10.0.0.36/30 [110/2] via 10.0.0.1, 02:28:56, FastEthernet0/1
C    10.0.0.40/30 is directly connected, FastEthernet1/0
L    10.0.0.41/32 is directly connected, FastEthernet1/0
O    10.0.0.44/30 [110/3] via 10.0.0.10, 02:28:46, FastEthernet0/0
O    10.0.0.48/30 [110/3] via 10.0.0.10, 02:28:46, FastEthernet0/0
Porto#

```

Figure 1.8: Porto router learned routes

```
Porto#show mpls forwarding-table
```

Local Label	Outgoing Label	Prefix or Tunnel Id	Bytes Switched	Label	Outgoing interface	Next Hop
16	Pop Label	10.0.0.36/30	0		Fa0/1	10.0.0.1
17	Pop Label	10.0.0.24/30	0		Fa0/1	10.0.0.1
18	Pop Label	10.0.0.20/30	0		Fa0/0	10.0.0.10
19	Pop Label	10.0.0.16/30	0		Fa0/0	10.0.0.10
20	Pop Label	10.0.0.12/30	0		Fa0/1	10.0.0.1
	Pop Label	10.0.0.12/30	0		Fa0/0	10.0.0.10
21	19	10.0.0.48/30	0		Fa0/0	10.0.0.10
22	20	10.0.0.44/30	0		Fa0/0	10.0.0.10
23	22	10.0.0.32/30	0		Fa0/0	10.0.0.10
24	24	10.0.0.4/30	0		Fa0/0	10.0.0.10

Figure 1.9: Porto router MPLS forwarding-table

```
10.0.0.0/8 is variably subnetted, 17 subnets, 2 masks
O 10.0.0.0/30 [110/3] via 10.0.0.22, 02:37:16, FastEthernet0/0
C 10.0.0.4/30 is directly connected, FastEthernet0/1
L 10.0.0.5/32 is directly connected, FastEthernet0/1
O 10.0.0.8/30 [110/2] via 10.0.0.22, 02:37:16, FastEthernet0/0
O 10.0.0.12/30 [110/2] via 10.0.0.22, 02:37:15, FastEthernet0/0
O 10.0.0.16/30 [110/2] via 10.0.0.22, 02:37:15, FastEthernet0/0
  [110/2] via 10.0.0.6, 02:37:15, FastEthernet0/1
C 10.0.0.20/30 is directly connected, FastEthernet0/0
L 10.0.0.21/32 is directly connected, FastEthernet0/0
O 10.0.0.24/30 [110/3] via 10.0.0.22, 02:37:16, FastEthernet0/0
O 10.0.0.28/30 [110/3] via 10.0.0.22, 02:37:16, FastEthernet0/0
C 10.0.0.32/30 is directly connected, FastEthernet1/1
L 10.0.0.33/32 is directly connected, FastEthernet1/1
O 10.0.0.36/30 [110/3] via 10.0.0.22, 02:37:16, FastEthernet0/0
O 10.0.0.40/30 [110/3] via 10.0.0.22, 02:37:16, FastEthernet0/0
O 10.0.0.44/30 [110/2] via 10.0.0.6, 02:37:16, FastEthernet0/1
C 10.0.0.48/30 is directly connected, FastEthernet1/0
L 10.0.0.49/32 is directly connected, FastEthernet1/0
Barcelona#
```

Figure 1.10: Barcelona router learned routes

```
Barcelona#show mpls forwarding-table
```

Local Label	Outgoing Label	Prefix or Tunnel Id	Bytes Switched	Label	Outgoing interface	Next Hop
16	Pop Label	10.0.0.44/30	0		Fa0/1	10.0.0.6
17	21	10.0.0.40/30	0		Fa0/0	10.0.0.22
18	16	10.0.0.36/30	0		Fa0/0	10.0.0.22
19	23	10.0.0.28/30	0		Fa0/0	10.0.0.22
20	17	10.0.0.24/30	0		Fa0/0	10.0.0.22
21	18	10.0.0.0/30	0		Fa0/0	10.0.0.22
22	Pop Label	10.0.0.8/30	0		Fa0/0	10.0.0.22
23	Pop Label	10.0.0.12/30	0		Fa0/0	10.0.0.22
24	Pop Label	10.0.0.16/30	0		Fa0/1	10.0.0.6
	Pop Label	10.0.0.16/30	0		Fa0/0	10.0.0.22

Figure 1.11: Barcelona router MPLS forwarding-table

Note: To obtain these tables we only turned on routers from Lisbon, Porto, Chicago, Barcelona and Core1, the rest of the network is off.

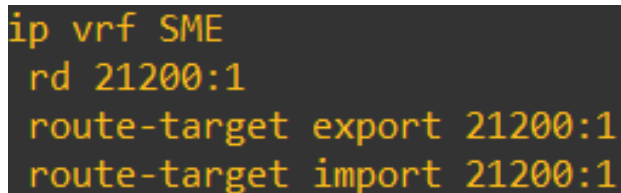
Chapter 2

SME1, SME2, and SME3 association private network

2.1 Implementation

For this part of the project we implemented a VRF which is a virtual routing and forwarding table for each router that connects to the different clients of the SME industrial association. In this way we can segment all the traffic that is generated by these clients.

All routers responsible for forwarding SME traffic have very similar configurations, instead of going through each one of them we will only show RL2 configuration.

A screenshot of a terminal window showing the configuration of a VRF named SME. The text is as follows:

```
ip vrf SME
rd 21200:1
route-target export 21200:1
route-target import 21200:1
```

Figure 2.1: SME VRF instance from RL2

This configuration sets up a VRF named SME with a unique identifier (RD 21200:1) and defines its routing policies:

- Routes originating from SME are tagged with RT 21200:1 when advertised (exported).
- The VRF will also accept and use routes tagged with RT 21200:1 (imported).

```

router bgp 21200
  bgp router-id 1.1.1.1
  bgp log-neighbor-changes
  neighbor 192.0.0.2 remote-as 21200
  neighbor 192.0.0.2 update-source Loopback0
  neighbor 192.0.0.3 remote-as 21200
  neighbor 192.0.0.3 update-source Loopback0
  neighbor 192.0.0.7 remote-as 21200
  neighbor 192.0.0.7 update-source Loopback0
  !
  address-family vpnv4
    neighbor 192.0.0.2 activate
    neighbor 192.0.0.2 send-community both
    neighbor 192.0.0.3 activate
    neighbor 192.0.0.3 send-community both
    neighbor 192.0.0.7 activate
    neighbor 192.0.0.7 send-community both
  exit-address-family
  !
  address-family ipv4 vrf SME
    redistribute connected
  exit-address-family

```

Figure 2.2: RL2 configuration

The SME VRF instance is exactly the same for all routers responsible to manage traffic generated from SME clients, what differs from router to router is mainly the BGP neighbors, since RL2 has the address 192.0.0.1 on its loopback interface, this router will create a BGP neighbor relation with all other routers that forward SME traffic, in this specific case RP2 router has the address 192.0.0.2, RB1 has the address 192.0.0.3 and the Barcelona router has the address 192.0.0.7. The routes are learned through BGP and OSPF protocols, MPLS LDP also gives an unique label in order to differentiate traffic generated by SME clients, as seen in the image below.

```

RL2#show ip bgp vpnv4 all labels

```

Network	Next Hop	In label/Out label
Route Distinguisher: 21200:1 (SME)		
10.0.0.48/30	192.0.0.7	nolabel/28
10.0.1.0/24	192.0.0.2	nolabel/30
10.0.2.0/24	0.0.0.0	30/nolabel(SME)
10.0.3.0/24	192.0.0.3	nolabel/30

Figure 2.3: Attributed label to SME clients

Performing a ping from client SME2 to SME1 we can see that the packet will

have 2 MPLS labels, this is because packets originated from the SME network will be given one label and when that packet enters the MPLS core it will be give another label depending on the link where the packet will go through.

No.	Time	Source	Destination	Protocol	Length	Info
33	21.920971	10.0.0.2	224.0.0.2	LDLP	76	Hello Message
34	21.920971	10.0.1.1	10.0.2.1	ICMP	122	Echo (ping) reply
35	21.937770	10.0.2.1	10.0.1.1	ICMP	122	Echo (ping) request
36	21.950950	10.0.1.1	10.0.2.1	ICMP	122	Echo (ping) reply
37	21.968619	10.0.2.1	10.0.1.1	ICMP	122	Echo (ping) request
38	21.982834	10.0.1.1	10.0.2.1	ICMP	122	Echo (ping) reply
39	21.999602	10.0.2.1	10.0.1.1	ICMP	122	Echo (ping) request
40	22.013743	ca:02:2f:40:00:06	ca:02:2f:40:00:06	LOOP	60	Reply
41	22.013743	10.0.1.1	10.0.2.1	ICMP	122	Echo (ping) reply

> Frame 37: 122 bytes on wire (976 bits), 122 bytes captured (976 bits) on interf
 > Ethernet II, Src: ca:01:21:a8:00:06 (ca:01:21:a8:00:06), Dst: ca:02:2f:40:00:06
 > MultiProtocol Label Switching Header, Label: 30, Exp: 0, S: 0, TTL: 254
 > MultiProtocol Label Switching Header, Label: 30, Exp: 0, S: 1, TTL: 255
 > Internet Protocol Version 4, Src: 10.0.2.1, Dst: 10.0.1.1
 > Internet Control Message Protocol

Figure 2.4: Packet capture that shows ping performed from SME2 to SME1

```
RL2#ping vrf SME 10.0.1.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.0.1.1, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 80/89/96 ms
```

Figure 2.5: SME2 to SME1 connectivity

```
RL2#show mpls forwarding-table
```

Local Label	Outgoing Label	Prefix or Tunnel Id	Bytes Switched	Outgoing interface	Next Hop
16	16	192.0.0.7/32	0	Fa0/0	10.0.0.25
17	21	10.0.0.44/30	0	Fa0/0	10.0.0.25
18	23	10.0.0.40/30	0	Fa0/0	10.0.0.25
19	Pop Label	10.0.0.36/30	0	Fa0/0	10.0.0.25
20	19	10.0.0.32/30	0	Fa0/0	10.0.0.25
21	20	10.0.0.28/30	0	Fa0/0	10.0.0.25
22	22	10.0.0.4/30	0	Fa0/0	10.0.0.25
23	25	10.0.0.20/30	0	Fa0/0	10.0.0.25
24	24	10.0.0.8/30	0	Fa0/0	10.0.0.25
25	26	10.0.0.16/30	0	Fa0/0	10.0.0.25
26	Pop Label	10.0.0.12/30	0	Fa0/0	10.0.0.25
27	Pop Label	10.0.0.0/30	0	Fa0/0	10.0.0.25
28	29	192.0.0.3/32	0	Fa0/0	10.0.0.25
29	30	192.0.0.2/32	0	Fa0/0	10.0.0.25
30	No Label	10.0.2.0/24[V]	570	aggregate/SME	

Figure 2.6: RL2 MPLS forwarding-table

Here are the VRF tables, which demonstrate that we can successfully ping between the different PoPs of the SME clients.

```
Gateway of last resort is 10.0.0.25 to network 0.0.0.0

S*    0.0.0.0/0 [1/0] via 10.0.0.25
      10.0.0.0/8 is variably subnetted, 5 subnets, 3 masks
B      10.0.0.48/30 [200/0] via 192.0.0.7, 00:19:10
B      10.0.1.0/24 [200/0] via 192.0.0.2, 00:19:04
C      10.0.2.0/24 is directly connected, FastEthernet0/1
L      10.0.2.1/32 is directly connected, FastEthernet0/1
B      10.0.3.0/24 [200/0] via 192.0.0.3, 00:19:05
RL2#
```

Figure 2.7: RL2 VRF table

```
Gateway of last resort is 10.0.0.29 to network 0.0.0.0

S*    0.0.0.0/0 [1/0] via 10.0.0.29
      10.0.0.0/8 is variably subnetted, 5 subnets, 3 masks
B      10.0.0.48/30 [200/0] via 192.0.0.7, 00:21:11
C      10.0.1.0/24 is directly connected, FastEthernet0/1
L      10.0.1.1/32 is directly connected, FastEthernet0/1
B      10.0.2.0/24 [200/0] via 192.0.0.1, 00:21:11
B      10.0.3.0/24 [200/0] via 192.0.0.3, 00:21:11
RP2#
```

Figure 2.8: RP2 VRF table

```
Gateway of last resort is 10.0.0.33 to network 0.0.0.0

S*    0.0.0.0/0 [1/0] via 10.0.0.33
      10.0.0.0/8 is variably subnetted, 5 subnets, 3 masks
B      10.0.0.48/30 [200/0] via 192.0.0.7, 00:23:39
B      10.0.1.0/24 [200/0] via 192.0.0.2, 00:23:39
B      10.0.2.0/24 [200/0] via 192.0.0.1, 00:23:39
C      10.0.3.0/24 is directly connected, FastEthernet0/1
L      10.0.3.1/32 is directly connected, FastEthernet0/1
RB1#
```

Figure 2.9: RB1 VRF table

Chapter 3

SME1, SME2, and SME3 CDN advanced service routing

3.1 Implementation

For this part of the project we implemented a conditional CDN service, that will answer to DNS queries based on the IP address of the client, for instance, clients SME1, SME2 and SME3 will have different service providers because these clients are on different networks, a server that is suitable for a client might not be suitable for the other, that is why it is important to define different "views" in the master server DNS service (which is located in the Barcelona PoP) in order to make the transmission of the service more efficient.

To start with, we defined a view in the file `"/etc/bind/named.conf.local"` of the master server for each SME client, in this way the bind server is able to handle the queries of each client in a specific way based on their IP address. If the client's network matches any view of the file, the server accesses that view where there is a database that says the IP address of the master server which answers to the queries made by clients, and there is also the IP address of the best service provider for the specific client that performs the query.


```

view "sme1" {
    match-clients {10.0.1.0/24};
    recursion no;
    zone "cdn4all.com" {
        type master;
        file "/etc/bind/cdn4all.com-sme1.db";
    };
};

view "sme2" {
    match-clients {10.0.2.0/24};
    recursion no;
    zone "cdn4all.com" {
        type master;
        file "/etc/bind/cdn4all.com-sme2.db";
    };
};

view "sme3" {
    match-clients {10.0.3.0/24};
    recursion no;
    zone "cdn4all.com" {
        type master;
        file "/etc/bind/cdn4all.com-sme3.db";
    };
};

```

Figure 3.1: Created views for each SME client

For example, the "sme1" view targets all clients within the 10.0.1.0/24 network. This view then accesses the database "cdn4all.com-sme1.db", which provides both the IP address of the master server and the most suitable service provider for the client.

The databases for each one of the SME clients are the following:

```

$TTL      604800
$ORIGIN    cdn4all.com.
@          IN      SOA      ns1.cdn4all.com. adm.cdn4all.com. (
                                2          ; Serial
                                604800     ; Refresh
                                86400      ; Retry
                                2419200; Expire
                                604800 ) ; Negative Cache TTL
@          IN      NS       ns1.cdn4all.com.
@          IN      A        10.0.1.2
ns1        IN      A        10.0.0.50

```

Figure 3.2: SME1 client database

- **Service Provider:** 10.0.1.2
- **Master Server:** 10.0.0.50

```

$TTL 604800
$ORIGIN cdn4all.com.
@      IN      SOA      ns1.cdn4all.com. adm.cdn4all.com. (
                                2          ; Serial
                                604800     ; Refresh
                                86400      ; Retry
                                2419200; Expire
                                604800 ) ; Negative Cache TTL
@      IN      NS       ns1.cdn4all.com.
@      IN      A        10.0.2.2
ns1    IN      A        10.0.0.50

```

Figure 3.3: SME2 client database

- **Service Provider:** 10.0.2.2
- **Master Server:** 10.0.0.50

```

$TTL 604800
$ORIGIN cdn4all.com.
@      IN      SOA      ns1.cdn4all.com. adm.cdn4all.com. (
                                2          ; Serial
                                604800     ; Refresh
                                86400      ; Retry
                                2419200; Expire
                                604800 ) ; Negative Cache TTL
@      IN      NS       ns1.cdn4all.com.
@      IN      A        10.0.3.2
ns1    IN      A        10.0.0.50

```

Figure 3.4: SME3 client database

- **Service Provider:** 10.0.3.2
- **Master Server:** 10.0.0.50

By performing a ping from the SME1, SME2 and SME3 networks to the "cdn4all.com" domain, we can confirm that the address is correctly translated, as indicated by the Wireshark captures.

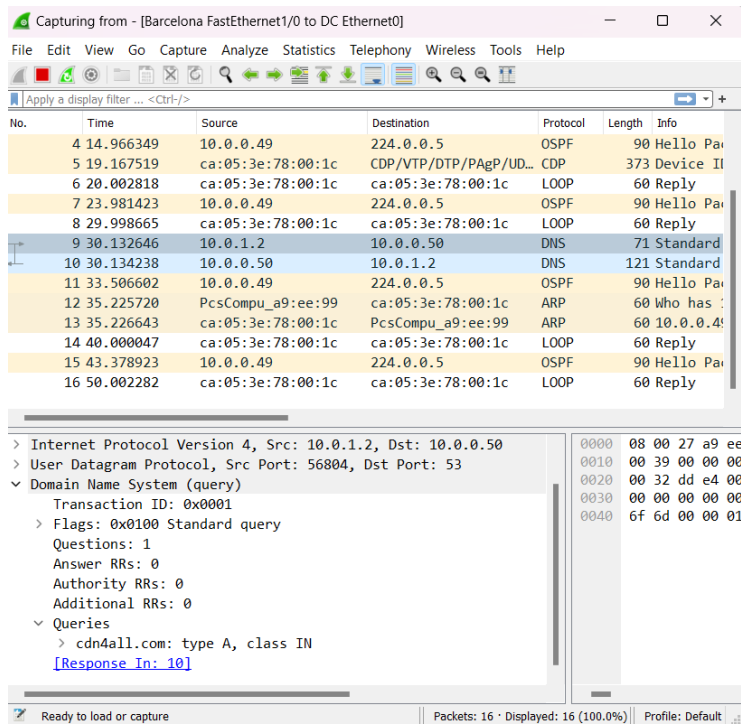


Figure 3.5: SME1 client query

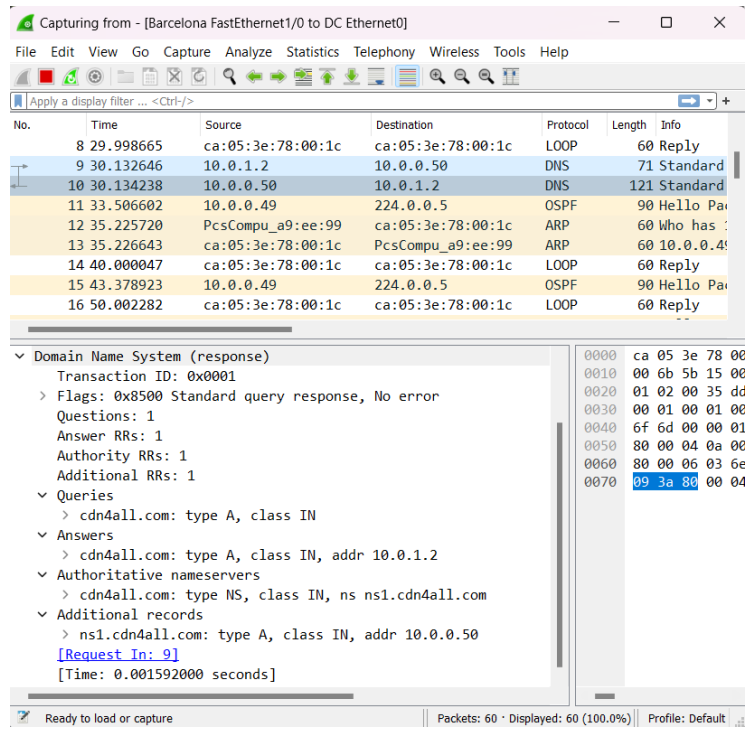


Figure 3.6: Master Server response

Note: As observed, the client requesting the service is the service provider itself. This approach was chosen because we decided to use the service provider as a client within the network instead of adding additional routers to represent clients. This decision was made to prevent significant performance degradation that could result from an excessive number of routers, which would greatly slow down our system.

```
Server1#ping cdn4all.com
Translating "cdn4all.com"...domain server (10.0.0.50) [OK]

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.0.1.2, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/3/8 ms
```

Figure 3.7: Successful ping from SME1 client to "cdn4all.com" domain

```
Server2#ping cdn4all.com
Translating "cdn4all.com"...domain server (10.0.0.50) [OK]

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.0.2.2, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/2/4 ms
```

Figure 3.8: Successful ping from SME2 client to "cdn4all.com" domain

```
Server3#ping cdn4all.com
Translating "cdn4all.com"...domain server (10.0.0.50) [OK]

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.0.3.2, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/5/8 ms
```

Figure 3.9: Successful ping from SME3 client to "cdn4all.com" domain

Chapter 4

Client LA Layer 2 private network

4.1 Assigned IPs to VPCs

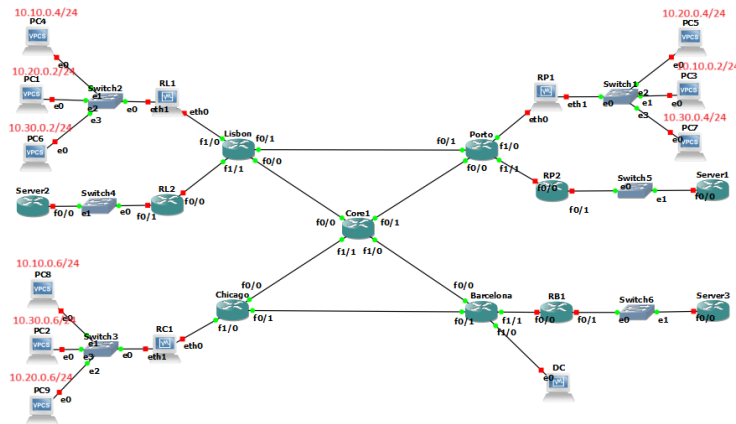


Figure 4.1: Network diagram with VPCs

4.2 Implementation

For this part of the project we utilized the protocol L2VPN/EVPN with VXLAN transport to obtain inter-connectivity across VLANs in different points of presence of the network. To segment the different VLANs we utilized cisco switches and VyOS routers.

The switch configuration is basically the same across all different PoPs:

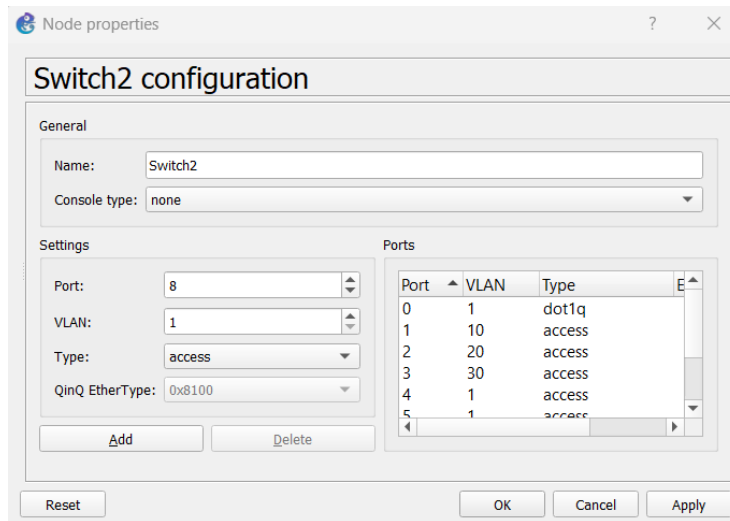


Figure 4.2: Switch2 configuration

We configured an access port for each VLAN (10, 20, and 30) to properly segment traffic originating from different VLANs. Additionally, we set up a trunk port connecting to the VyOS router, which allows traffic from all VLANs to pass through. This configuration enables routing of VLAN traffic to other Points of Presence (PoPs).

Moreover, we configured a bridge interface on the VyOS router for each VLAN, enabling logical network segmentation with its own MAC address space to handle Layer 2 forwarding decisions. To match traffic from each VLAN, we created three virtual interfaces and assigned a unique VXLAN Network Identifier (VNI) to each VXLAN. This setup allows us to distinguish between virtualized Layer 2 networks while utilizing a shared Layer 3 infrastructure.

```
shvyos@RC1:~$ show conf
interfaces {
    bridge br1010 {
        address 10.10.0.5/24
        description "vlan 10"
        member {
            interface eth1.10 {
            }
            interface vxlan1010 {
            }
        }
    }
    bridge br1020 {
        address 10.20.0.5/24
        description "vlan 20"
        member {
            interface eth1.20 {
            }
            interface vxlan1020 {
            }
        }
    }
}
```

Figure 4.3: RC1 Bridge 1010 and Bridge 1020


```

bridge br1030 {
    address 10.30.0.5/24
    description "vlan 30"
    member {
        interface eth1.30 {
        }
        interface vxlan1030 {
        }
    }
}
dummy dum0 {
    address 192.0.0.6/32
}
ethernet eth0 {
    address 10.0.0.46/30
    hw-id 08:00:27:a5:c9:3b
}
ethernet eth1 {
    hw-id 08:00:27:ae:27:ea
    vif 10 {
    }
    vif 20 {
    }
    vif 30 {
    }
}

```

Figure 4.4: RC1 Bridge 1030 and VIFs

```

vxlan vxlan1010 {
    mtu 1500
    source-address 192.0.0.6
    vni 1010
}
vxlan vxlan1020 {
    mtu 1500
    source-address 192.0.0.6
    vni 1020
}
vxlan vxlan1030 {
    mtu 1500
    source-address 192.0.0.6
    vni 1030
}

```

Figure 4.5: RC1 VXLANs configurations

This interconnection of the different VLAN PoPs is possible because of the established BGP mesh with other VyOS routers, for BGP peering relations, we used the Loopback interfaces of each router since they are less likely to go down, we assigned the IP **192.0.0.6** to router **RC1**, **192.0.0.5** to router **RP1** and IP **192.0.0.4** to router **RL1**. The VNIs will then be advertised through BGP allowing other PoPs to be aware of the presence of the same VLANs in different networks. In this way we can use the protocol EVPN to transport Layer 2 traffic over a Layer 3 infrastructure.

```

protocols {
    bgp {
        address-family {
            l2vpn-evpn {
                advertise-all-vni
            }
        }
        neighbor 192.0.0.4 {
            peer-group evpn
        }
        neighbor 192.0.0.5 {
            peer-group evpn
        }
        parameters {

```

Figure 4.6: RC1 BGP neighbors

```

parameters {
    router-id 192.0.0.6
}
peer-group evpn {
    address-family {
        l2vpn-evpn {
            nexthop-self {
            }
        }
    }
    remote-as 21200
    update-source dum0
}
system-as 21200
}

```

Figure 4.7: RC1 L2VPN/EVPN

```

ospf {
    area 0 {
        network 10.0.0.44/30
        network 192.0.0.6/32
    }
}

```

Figure 4.8: RC1 OSPF Process

The IP **10.0.0.44** corresponds to the network between the RC1 VyOS router and the Chicago router, the **192.0.0.6** IP is the address of the Loopback interface of the router.

Note: These configurations were applied to the RC1 router. The configurations on the other VyOS routers are nearly identical, with the only differences being the IP addresses of the Loopback interfaces and the BGP neighbor settings.

When performing a ping from PC4 (10.10.0.4) to PC8 (10.10.0.6), the ping is successful. Additionally, the Wireshark capture confirms that the VNI is being correctly assigned to VXLAN 1010.

```
PC4> ping 10.10.0.6
84 bytes from 10.10.0.6 icmp_seq=1 ttl=64 time=52.101 ms
84 bytes from 10.10.0.6 icmp_seq=2 ttl=64 time=35.036 ms
84 bytes from 10.10.0.6 icmp_seq=3 ttl=64 time=45.423 ms
84 bytes from 10.10.0.6 icmp_seq=4 ttl=64 time=57.013 ms
84 bytes from 10.10.0.6 icmp_seq=5 ttl=64 time=56.744 ms
```

Figure 4.9: PC4 to PC8 ping

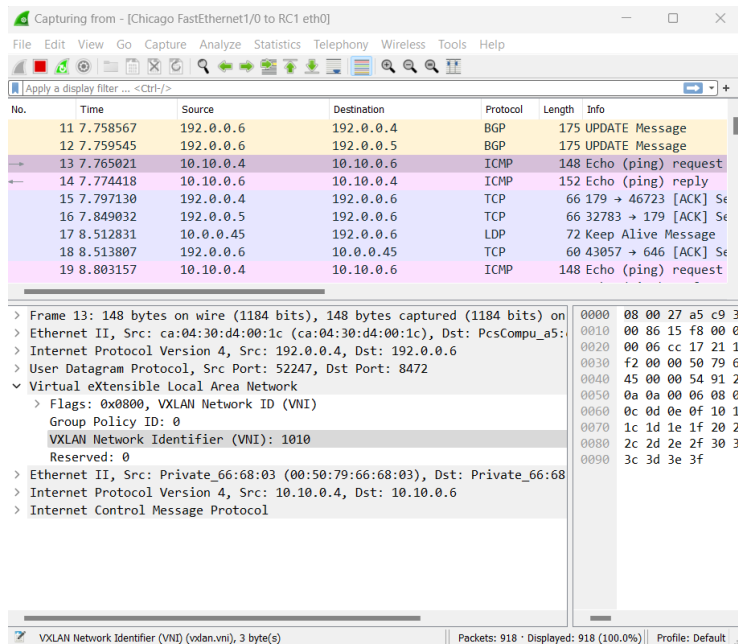


Figure 4.10: Wireshark capture from VLAN 10 ping

Performing a ping from PC1 (10.20.0.2) to PC9 (10.20.0.6) also shows there is connection, and VXLAN 1020 is correctly identified.

```

PC1> ping 10.20.0.6
84 bytes from 10.20.0.6 icmp_seq=1 ttl=64 time=62.110 ms
84 bytes from 10.20.0.6 icmp_seq=2 ttl=64 time=61.194 ms
84 bytes from 10.20.0.6 icmp_seq=3 ttl=64 time=44.646 ms
84 bytes from 10.20.0.6 icmp_seq=4 ttl=64 time=51.807 ms
84 bytes from 10.20.0.6 icmp_seq=5 ttl=64 time=51.890 ms

```

Figure 4.11: PC1 to PC9 ping

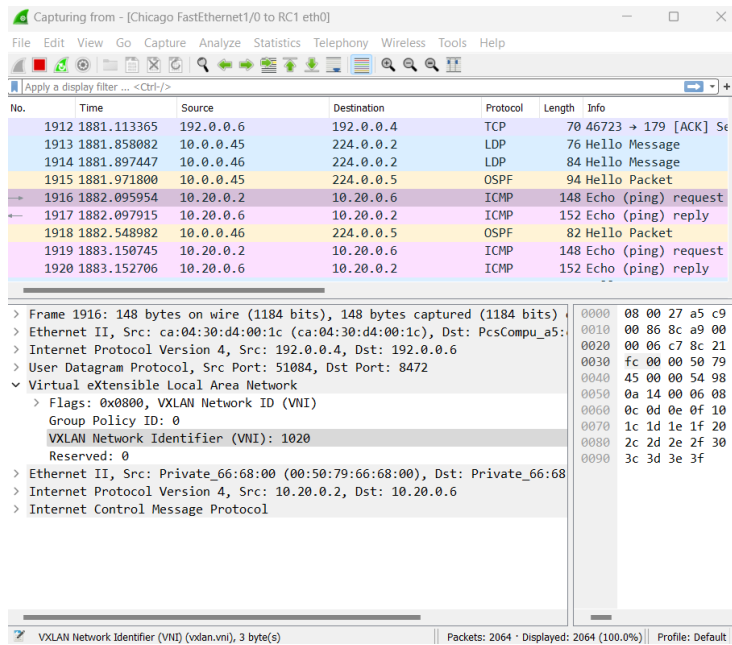


Figure 4.12: Wireshark capture from VLAN 20 ping

Performing a ping from PC6 (10.30.0.2) to PC2 (10.30.0.6) also yields the same results, but with a different VNI.

```

PC6> ping 10.30.0.6
84 bytes from 10.30.0.6 icmp_seq=1 ttl=64 time=50.384 ms
84 bytes from 10.30.0.6 icmp_seq=2 ttl=64 time=51.676 ms
84 bytes from 10.30.0.6 icmp_seq=3 ttl=64 time=51.718 ms
84 bytes from 10.30.0.6 icmp_seq=4 ttl=64 time=51.900 ms
84 bytes from 10.30.0.6 icmp_seq=5 ttl=64 time=52.827 ms

```

Figure 4.13: PC6 to PC2 ping

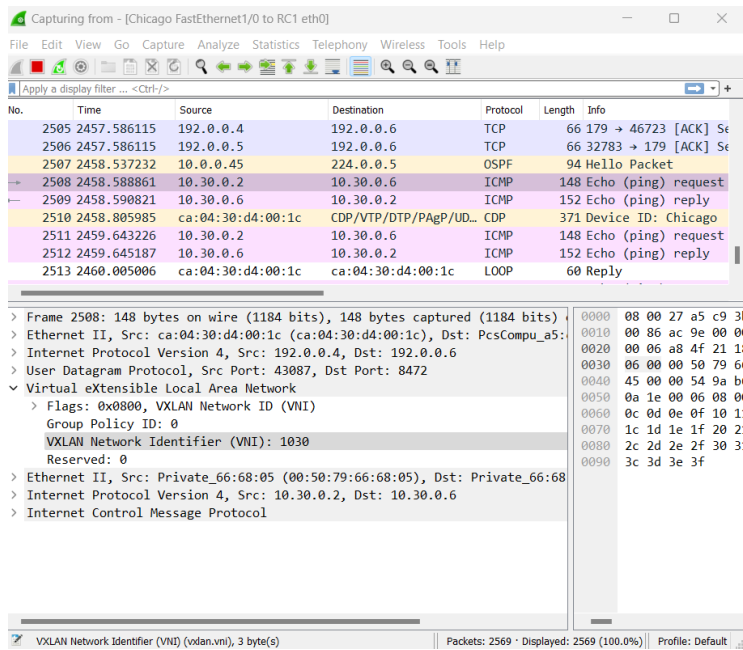


Figure 4.14: Wireshark capture from VLAN 30 ping

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 192.0.0.4:2					
*>i[3]:[0]:[32]:[192.0.0.4]	192.0.0.4		100	0	i
	RT:21200:1030 ET:8				
Route Distinguisher: 192.0.0.4:3					
*>i[3]:[0]:[32]:[192.0.0.4]	192.0.0.4		100	0	i
	RT:21200:1010 ET:8				
Route Distinguisher: 192.0.0.4:4					
*>i[3]:[0]:[32]:[192.0.0.4]	192.0.0.4		100	0	i
	RT:21200:1020 ET:8				
Route Distinguisher: 192.0.0.5:2					
*>i[3]:[0]:[32]:[192.0.0.5]	192.0.0.5		100	0	i
	RT:21200:1030 ET:8				

Figure 4.15: RC1 EVPN prefixes part 1

```

Route Distinguisher: 192.0.0.5:3
*>i[31:[0]:[32]:[192.0.0.5]
      192.0.0.5                100      0 i
      RT:21200:1010 ET:8
Route Distinguisher: 192.0.0.5:4
*>i[31:[0]:[32]:[192.0.0.5]
      192.0.0.5                100      0 i
      RT:21200:1020 ET:8
Route Distinguisher: 192.0.0.6:2
*> [31:[0]:[32]:[192.0.0.6]
      192.0.0.6                32768 i
      ET:8 RT:21200:1030
Route Distinguisher: 192.0.0.6:3
*> [31:[0]:[32]:[192.0.0.6]
      192.0.0.6                32768 i
      ET:8 RT:21200:1010
Route Distinguisher: 192.0.0.6:4
*> [31:[0]:[32]:[192.0.0.6]
      192.0.0.6                32768 i
      ET:8 RT:21200:1020
Displayed 9 out of 9 total prefixes

```

Figure 4.16: RC1 EVPN prefixes part 2

Note: The pings were performed only between the Lisbon and Chicago PoPs. Including the results of pings to the Porto PoP would be redundant, as the connection between the PoPs can be easily demonstrated during the project presentation.

Chapter 5

Client LA Layer 2 traffic differentiation

This part of the project aims to differentiate traffic (at the core of the network) generated by client LA by defining an Assured Forwarding policy, that should be guaranteed up to 10 Mbps.

We applied several configurations on routers from Lisbon, Porto, Chicago and Core1, but we will only cover the configurations made on routers from Lisbon and Core1.

```
class-map match-all GOLD
  match ip dscp af11
class-map match-all AF11
  match access-group 101
!
policy-map SETDSCP
  class AF11
    set ip dscp af11
policy-map EDGE
  class GOLD
    bandwidth 10000
```

Figure 5.1: Lisbon router class and policy maps


```
access-list 101 permit udp host 192.0.0.4 host 192.0.0.5 eq 8472
access-list 101 permit udp host 192.0.0.4 host 192.0.0.6 eq 8472
```

Figure 5.2: Access-list 101 from Lisbon router

```
interface FastEthernet0/0
 ip address 10.0.0.13 255.255.255.252
 ip ospf 1 area 0
 speed auto
 duplex auto
 mpls ip
 service-policy output EDGE
!
interface FastEthernet0/1
 ip address 10.0.0.1 255.255.255.252
 ip ospf 1 area 0
 speed auto
 duplex auto
 mpls ip
 service-policy output EDGE
!
interface FastEthernet1/0
 ip address 10.0.0.37 255.255.255.252
 ip ospf 1 area 0
 speed auto
 duplex auto
 service-policy input SETDSCP
```

Figure 5.3: Lisbon router interfaces

Basically, all traffic that will be forwarded to port 8472, will belong to the GOLD class which uses assured forwarding (AF11) with low drop probability (dscp 10). The DSCP value should be attributed to the packet in the input interface of the router which in this case is "FastEthernet1/0" and the policy is enforced in the output interfaces which can be either "FastEthernet0/0" or "FastEthernet0/1".

We will exclude the Porto and Chicago routers from our detailed configuration, as their setups are identical to the Lisbon router. The only difference lies in the access control lists (ACLs) defined on each router. For example, the Porto router has two ACLs: one to filter traffic from Porto to Lisbon and another to filter traffic from Porto to Chicago. Similarly, the Chicago router has two ACLs: one for traffic from Chicago to Porto and another for traffic from Chicago to Lisbon.

Moreover, Core1 plays a crucial role in extending the scope of QoS for client LA. This router ensures that QoS continues to be enforced across the network by identifying packets belonging to the GOLD class and allocating up to 10 Mbps on the output interfaces for client LA traffic.

```
interface FastEthernet0/0
ip address 10.0.0.14 255.255.255.252
ip ospf 1 area 0
speed auto
duplex auto
mpls ip
service-policy output EDGE
!
interface FastEthernet0/1
ip address 10.0.0.10 255.255.255.252
ip ospf 1 area 0
speed auto
duplex auto
mpls ip
service-policy output EDGE
!
interface FastEthernet1/0
ip address 10.0.0.22 255.255.255.252
ip ospf 1 area 0
speed auto
duplex auto
mpls ip
service-policy output EDGE
!
interface FastEthernet1/1
ip address 10.0.0.18 255.255.255.252
ip ospf 1 area 0
speed auto
duplex auto
mpls ip
service-policy output EDGE
```

Figure 5.4: Core1 router interfaces

```

class-map match-all GOLD
  match ip dscp af11
!
policy-map EDGE
  class GOLD
    bandwidth 10000

```

Figure 5.5: Core1 router class and policy maps

In conclusion, the following screenshot verifies that the "diffserv" functionality has been correctly implemented for client LA. To test this, we initiated a ping from PC1 (10.20.0.2) to PC9 (10.20.0.6) and captured the traffic using Wireshark on the link between Chicago and Core1.

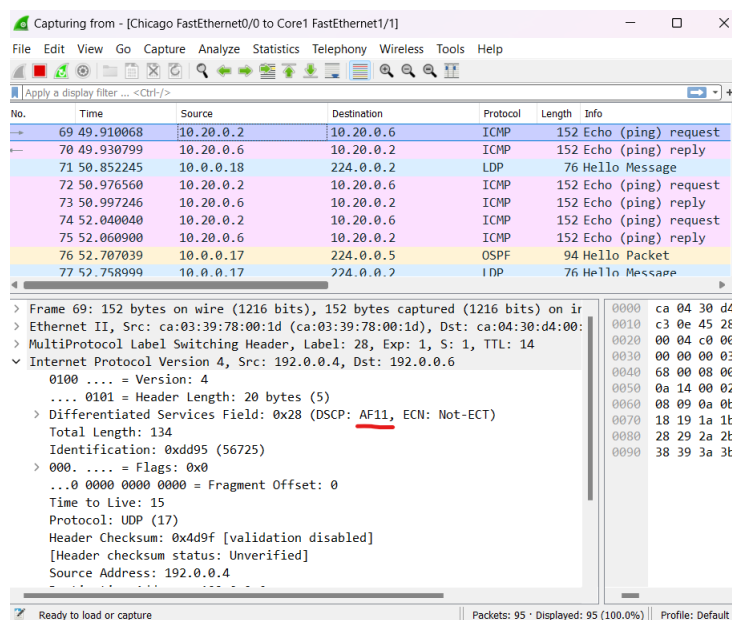


Figure 5.6: Performed ping from PC1 to PC9

Important Note: During the project demonstration with Professor Paulo Salvador, we had not yet implemented this functionality correctly. However, Professor Paulo provided valuable advice on how to properly implement it. After the demonstration, we successfully applied the necessary changes to ensure the

feature was correctly implemented. While we are unsure if this will positively impact our grade, we would greatly appreciate it if it did.

Chapter 6

SME1, SME2, and SME3 private networks traffic reserves

Not implemented.

Chapter 7

Conclusion

In this project, we successfully designed, configured, and tested a CDN network infrastructure capable of meeting the requirements of multiple enterprise clients. Key implementations included the setup of private networks for SME clients, advanced service routing for CDN servers, the deployment of a Layer 2 private network for Client LA and traffic differentiation.

However, due to time management issues, the implementation of traffic reservation for SME clients was not implemented. This feature remains as an opportunity for future development to further enhance the network's performance and client-specific quality of service (QoS).

Overall, the project demonstrates a scalable and robust approach to managing multi-client network architectures while addressing diverse connectivity requirements.

Author's Contribution

Each member contributed equally to the development of this project, therefore, we decided to allocate 50% to each.