

A black spiral-bound notebook is shown from a top-down perspective, lying flat on a wooden surface. The notebook's cover is a solid black. A teal-colored banner with a slightly textured, paper-like appearance is centered on the cover. The banner has a horizontal rectangular shape with pointed ends on the left and right sides. On the banner, the text "How does work" is written in a white, casual, handwritten-style font on the top line, and "JPEG?" is written on the bottom line. To the right of the banner, a piece of light-colored, braided twine is wrapped around the notebook's spine and tied into a simple knot. The left edge of the notebook shows the silver-colored metal spiral binding. The background is a dark brown wooden surface with visible wood grain patterns.


How does work
JPEG?

JPEG compression

- JPEG Compression is lossy, such that the original image cannot be exactly reconstructed (although a "Lossless JPEG" specification exists as well).
- JPEG exploits the characteristics of human vision, eliminating or reducing data to which the eye is less sensitive. JPEG works well on grayscale and color images, especially on photographs, but is not intended for two-tone images.



original: 104 KB





compressed: 92 KB (-11%)



crompressed: 5 KB (-95%)

Idea behind compression (JPEG)

- Our eyes are highly sensitive to intensity (grayscale) information and less sensitive to color information.
- Again we are more sensitive to low frequency intensity value and less sensitive to high frequency intensity value.

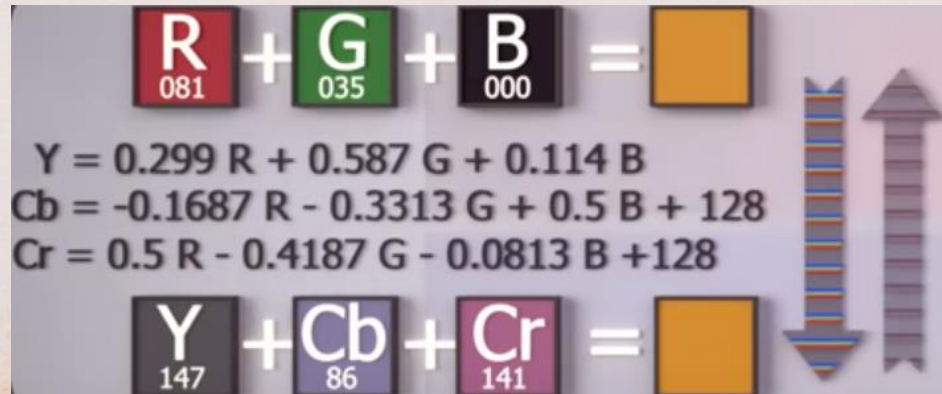


Steps of compression

- 01 color Space conversion
- 02 chrominance Downsampling
- 03 Discrete cosine Transformation
- 04 quantization
- 05 Run Length , Huffman Encoding and Arithmetic encoding

color Space conversion

- Calculates 3 new luminance values:
 - Blue chrominance
 - Red chrominance
 - Y luminance
- This process is reversible and there is no loss of data



The diagram illustrates the conversion from RGB to YCbCr color space. At the top, three colored squares represent the input: a red square labeled 'R' with the value '081' below it, a green square labeled 'G' with '035' below it, and a blue square labeled 'B' with '000' below it. These are followed by a plus sign and an equals sign, leading to an orange square. Below this, the conversion formulas are listed:
$$Y = 0.299 R + 0.587 G + 0.114 B$$
$$Cb = -0.1687 R - 0.3313 G + 0.5 B + 128$$
$$Cr = 0.5 R - 0.4187 G - 0.0813 B + 128$$
At the bottom, three colored squares represent the output: a yellow square labeled 'Y' with the value '147' below it, a blue square labeled 'Cb' with '86' below it, and a magenta square labeled 'Cr' with '141' below it. These are followed by a plus sign and an equals sign, leading to another orange square. To the right of the equations, there are two vertical arrows: a downward-pointing arrow with horizontal stripes and an upward-pointing arrow with horizontal stripes, indicating the reversible nature of the conversion.




$R_{081} + G_{035} + B_{000} = \text{Orange}$

$$Y = 0.299 R + 0.587 G + 0.114 B$$
$$Cb = -0.1687 R - 0.3313 G + 0.5 B + 128$$
$$Cr = 0.5 R - 0.4187 G - 0.0813 B + 128$$

$Y_{147} + Cb_{86} + Cr_{141} = \text{Orange}$

chrominance Downsampling

- In the next step we remove a considerable amount of data, remember that our eyes are bad at detecting color or prominence versus brightness or luminance.
- In general, jpeg down sample by 2 in each direction. Means that the color information is down sampled by 4 (4 time less color than the original one).



Before: $1+1+1=3.0$

After: $1/4+1/4+1=1.5$

Preprocessing Discrete cosine Transformation


- First the image is separated into non-overlapping blocks where each block is 8 X 8 pixels(64 pixels per block).
- Then normalize the pixels value to -128 to 127 from 0 to 255, just by subtracting 128 from each pixel value.

037 -128 -091	036 -128 -092	037 -128 -091	042 -128 -086	047 -128 -081	050 -128 -078	055 -128 -073	061 -128 -067
037 -128 -091	037 -128 -091	038 -128 -090	044 -128 -084	049 -128 -079	052 -128 -076	054 -128 -074	057 -128 -071
038 -128 -090	039 -128 -089	044 -128 -084	052 -128 -076	056 -128 -072	057 -128 -071	056 -128 -072	056 -128 -072
037 -128 -091	044 -128 -084	056 -128 -072	065 -128 -063	068 -128 -060	067 -128 -061	062 -128 -066	059 -128 -069
043 -128 -085	057 -128 -071	072 -128 -056	083 -128 -045	085 -128 -043	080 -128 -048	073 -128 -055	066 -128 -062
054 -128 -074	057 -128 -071	097 -128 -031	105 -128 -023	102 -128 -026	093 -128 -035	084 -128 -044	077 -128 -051
072 -128 -056	101 -128 -027	117 -128 -011	116 -128 -013	111 -128 -017	105 -128 -023	095 -128 -033	085 -128 -043
091 -128 -037	115 -128 -013	124 -128 -008	121 -128 -011	116 -128 -016	108 -128 -024	100 -128 -028	092 -128 -036

Discrete cosine Transformation

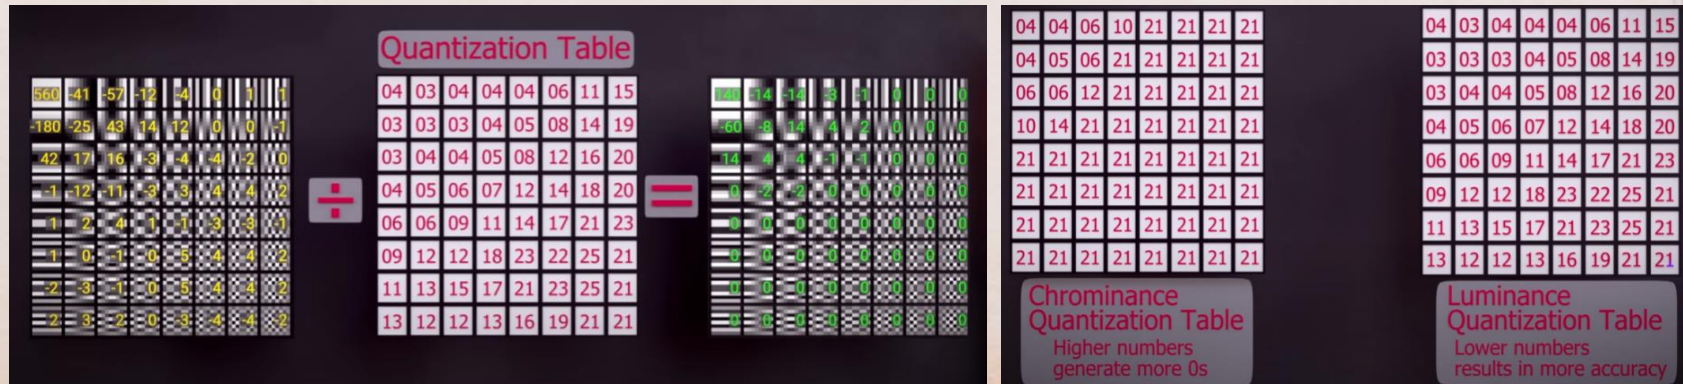
- DCT creates a new representation of image where it is easy to separate important and less important pixels of image.
- Each "grid" or pattern corresponds to a specific type of variation:
The first grid (in the top-left corner) represents uniform areas, meaning very low frequencies.
Grids further down and to the right represent faster and more complex variations, meaning higher frequencies.
- Multiplication by coefficients:
Each of these grids has a weight or coefficient that indicates how important that frequency pattern is to the original block.

What happens is that each pattern (grid) is multiplied by its coefficient and then summed up to reconstruct the image.



Quantization


- This is where the data is compressed in great. After performing DCT, we apply quantization on it to remove the high frequency components. This is the real data comprehension part of jpeg.
- For quantizing the DCT transformed data, jpeg use a predefined quantization table. The amount of comprehension is dependent on the choice of quantization table.
- This quantization is non reversible meaning that we can't recover the original data without loss.



Zig Zag Scanning

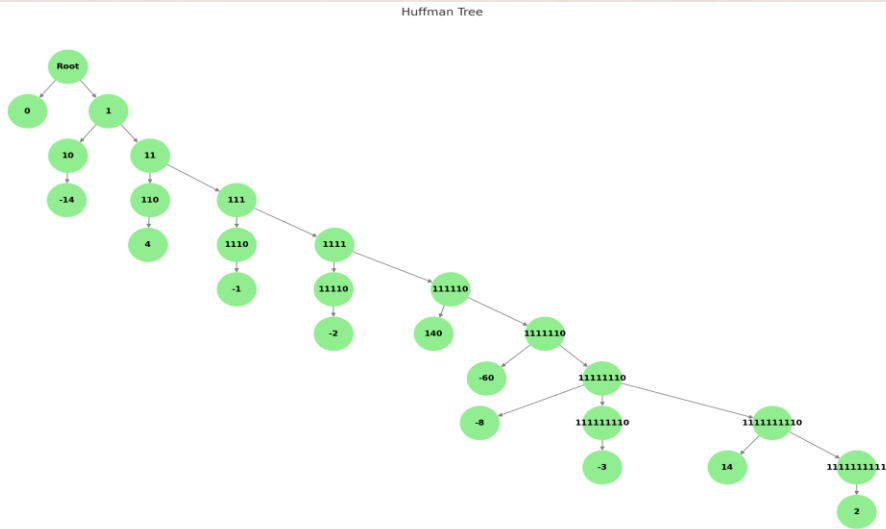
- The zigzag ordering groups non-zero coefficients (important information) at the beginning of the sequence and pushes all the zeros (less important information) to the end.
- This makes it easier to apply further compression methods, like Run-Length Encoding (RLE), which can represent long sequences of zeros efficiently.
- The zigzag path (highlighted in the image) starts at the top-left (low frequencies) and gradually moves toward the bottom-right (high frequencies).
- Run Length Encoding

140, -14, -60, -14, -8, -14, -3, 14, 4, 0[x2], -2, 4[x2], -1, 0, 2, -1,
-2, 0[x6], -1, 0[x38]



Huffman coding

140, -14, -60, -14, -8, -14, -3, 14, 4, 0, 0, -2, 4, 4,
-1, 0, 2, -1, -2, 0, 0, 0, 0, 0, 0, -1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0



Value	Frequency
0	40
-14	3
4	3
-1	2
-2	2
140	1
-60	1
-8	1
-3	1
14	1
2	1

Huffman coding

140, -14, -60, -14, -8, -14, -3, 14, 4, 0, 0, -2, 4, 4, -1, 0,
2, -1, -2, 0, 0, 0, 0, 0, 0, -1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0

|||||0 10 |||||0 10 |||||0 10 |||||0 10 |||||0 10

110 0 0 1110 110 110 1110 0 11111111 1110 1110

0 0 0 0 0 0 1110 0 0 0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0

original: 50 values * 8 bits = 400 bits

compressed: $(1 \times 6) + (3 \times 2) + (1 \times 7) + (1 \times 8) + (1 \times 9) + (1 \times 10) +$
 $(3 \times 3) + (40 \times 1) + (2 \times 5) + (2 \times 4) + (1 \times 10) = 123 \text{ bits}$

Value	Huffman Code
0	0
-14	10
4	110
-1	1110
-2	11110
140	111110
-60	1111110
-8	11111110
-3	111111110
14	1111111110
2	1111111111

Arithmetic coding

Value	Frequency	Probability
0	40	$40/50 = 0.8$
-14	3	$3/50 = 0.06$
4	3	$3/50 = 0.06$
-1	2	$2/50 = 0.04$
-2	2	$2/50 = 0.04$
140	1	$1/50 = 0.02$
-60	1	$1/50 = 0.02$
-8	1	$1/50 = 0.02$
-3	1	$1/50 = 0.02$
14	1	$1/50 = 0.02$
2	1	$1/50 = 0.02$

Value	Cumulative range
0	[0.00, 0.80)
-14	[0.80, 0.86)
4	[0.86, 0.92)
-1	[0.92, 0.96)
-2	[0.96, 1.00)
140	[0.00, 0.02)
-60	[0.02, 0.04)
-8	[0.04, 0.06)
-3	[0.06, 0.08)
14	[0.08, 0.10)
2	[0.10, 0.12)

Arithmetic coding

- Instead of assigning fixed-length codes to symbols, arithmetic coding represents an entire message as a single number between 0 and 1. It continuously subdivides the range $[0, 1)$ based on symbol probabilities.

First Few Symbols: 140, -14, -60,

Step 1: Start with $[0.00, 1.00)$

Step 2: Encode 140: Range of 140: $[0.00, 0.02)$

New range = $[0.00 + 0.00 \times (1.00 - 0.00), 0.00 + 0.02 \times (1.00 - 0.00)] = [0.00, 0.02)$

Step 3: Encode -14: Range of -14: $[0.80, 0.86)$

New range = $[0.00 + 0.80 \times (0.02 - 0.00), 0.00 + 0.86 \times (0.02 - 0.00)] = [0.016, 0.0172)$

Step 4: Encode -60: Range of -60: $[0.02, 0.04)$

New range = $[0.016 + 0.02 \times (0.0172 - 0.016), 0.016 + 0.04 \times (0.0172 - 0.016)] = [0.01604, 0.01608)$

JPEG 2000 VS JPEG

Feature	JPEG 2000	JPEG
Compression Quality	Higher, fewer artifacts	Lower, with noticeable blockiness
Compression Modes	Lossless & Lossy	Lossy only
Progressive Decoding	Yes	Yes (but less efficient)
Error Resilience	Better	Poor
Bit Depth Support	Up to 16 bits	8 bits
File Size	Smaller for same quality	Larger for same quality

Applications of JPEG in Digital Photography and Web Images

Digital Photography:

- High compression Efficiency;
- compatibility;
- Storage on Memory cards;
- Quick Sharing;
- common use case;



Web images:

- Efficient Web Page Loading;
- High-Quality Visuals;
- universal Browser Support;
- Responsive and Scalable Design;
- Typical use cases;

